

Day 1 Detailed Lab Implementation Guide

Prerequisites

1. **Local Machine Setup**

- Install **kubect** (v1.23+)

Detailed Guide: Installing kubect (v1.23+)

This guide provides platform-specific instructions to install **kubect**, the Kubernetes command-line tool, version 1.23 or higher.

1. Verify Prerequisites

- A supported operating system: macOS, Linux (Debian/Ubuntu, CentOS/RHEL, Fedora), or Windows.
- Internet access to download binaries.
- Sufficient privileges to install system packages or place binaries.

2. Install on Linux

2.1 Using Package Manager (Debian/Ubuntu)

1. Update package index and install dependencies:

```
sudo apt-get update
```

```
sudo apt-get install -y ca-certificates curl
```

2. Download v1.23 binary:

```
curl -LO "https://dl.k8s.io/release/v1.23.17/bin/linux/amd64/kubect"
```

3. Validate checksum:

```
curl -LO "https://dl.k8s.io/v1.23.17/bin/linux/amd64/kubect.sha256"
```

```
echo "$(cat kubect.sha256) kubect" | sha256sum --check
```

4. Install:

```
chmod +x kubect
```

```
sudo mv kubect /usr/local/bin/
```

5. Confirm:

```
kubectl version --client --short
```

3. Post-Installation Configuration

1. Autocomplete (optional):

- macOS/Linux (bash):

```
source <(kubectl completion bash)
```

- macOS/Linux (zsh):

```
source <(kubectl completion zsh)
```

- Windows (PowerShell):

```
kubectl completion powershell | Out-String | Invoke-Expression
```

2. Verify connectivity once you configure a cluster context:

```
kubectl cluster-info
```

kubectl is now installed and ready for use with Kubernetes v1.23 or newer.

- Install **Helm** (v3+)

Detailed Guide: Installing Helm (v3+)

This guide provides platform-specific instructions to install **Helm**, the Kubernetes package manager, version 3 or higher.

1. Verify Prerequisites

- A supported operating system: macOS, Linux (Debian/Ubuntu, CentOS/RHEL, Fedora), or Windows.
- Internet access to download binaries.
- Sufficient privileges to install system packages or place binaries.

2. Install on Linux

2.1 Using Package Manager (Debian/Ubuntu)

1. Update package index and install dependencies:

```
sudo apt-get update
```

```
sudo apt-get install -y apt-transport-https ca-certificates curl
```

2. Add Helm GPG key and repository:

```
curl https://baltocdn.com/helm/signing.asc | sudo apt-key add -
```

```
sudo apt-get install -y software-properties-common
```

```
sudo add-apt-repository "deb https://baltocdn.com/helm/stable/debian/ all main"
```

3. Install Helm:

```
sudo apt-get update
```

```
sudo apt-get install -y helm
```

4. Verify version:

```
helm version --short
```

3. Post-Installation Configuration

1. **Add Stable Repository** (optional but recommended):

```
helm repo add stable https://charts.helm.sh/stable
```

```
helm repo update
```

2. **Enable Autocompletion:**

- **bash:**

```
source <(helm completion bash)
```

- **zsh:**

```
source <(helm completion zsh)
```

3. **Verify Connectivity:**

Deploy a test chart (e.g., nginx) to confirm Helm can install releases:

```
helm install test-nginx stable/nginx-ingress --namespace default
```

```
helm list
```

Helm is now installed and ready for managing Kubernetes packages with version 3 or higher.

- Install ****Docker**** and login to a public registry (optional)

Detailed Guide: Installing Docker & Logging In to a Public Registry

This guide provides **step-by-step instructions** to install Docker on macOS, Linux, and Windows, and authenticate with Docker Hub (or any public registry).

1. Verify Prerequisites

- Supported OS: macOS, Linux (Debian/Ubuntu, CentOS/RHEL, Fedora), Windows 10/11.
- Internet connectivity.
- Sufficient privileges to install applications or run elevated commands.

2. Install Docker on Linux

2.1 Debian/Ubuntu

1. Update package index and install dependencies:

```
sudo apt-get update
```

```
sudo apt-get install -y apt-transport-https ca-certificates curl gnupg lsb-release
```

2. Add Docker's official GPG key:

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg
```

3. Add Docker repository:

```
echo \
```

```
"deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg] \
```

```
https://download.docker.com/linux/ubuntu \
```

```
$(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

4. Install Docker Engine:

```
sudo apt-get update
```

```
sudo apt-get install -y docker-ce docker-ce-cli containerd.io
```

5. Add your user to the Docker group (optional, for non-root usage):

```
sudo usermod -aG docker $USER
```

newgrp docker

6. Verify:

docker version

3. Log In to a Public Registry (Docker Hub)

1. Open a terminal, PowerShell, or Docker CLI.
2. Execute login command:

docker login

3. When prompted, enter your Docker Hub **username** and **password**.
4. On success, you will see:

Login Succeeded

5. (Optional) To log in to another registry, specify its URL:

docker login myregistry.example.com

6. Push & Pull a Test Image

1. Pull an official image:

docker pull alpine:latest

2. Tag it for your namespace:

docker tag alpine:latest <your-username>/alpine:test

3. Push to Docker Hub:

docker push <your-username>/alpine:test

4. Verify by pulling it again:

docker pull <your-username>/alpine:test

Docker and registry authentication are now configured, enabling you to build, push, and pull images seamlessly.

- Install ****Zeebe CLI**** (`zbctl`): ``brew install zeebe/zeebe/zbctl`` or download from <https://github.com/camunda-cloud/zeebe/releases>

- Ensure access to a Kubernetes cluster (e.g., Minikube, kind, or a managed cluster)

```
kubectl cluster-info
```

```
kubectl get nodes
```

```
kubectl get pods --all-namespaces
```

2. **Public Repositories**

- **Camunda 8 Helm Charts**: <https://github.com/camunda/camunda-platform-deployment>

- **Sample BPMN & Worker Project**: <https://github.com/camunda-community-hub/camunda-8-get-started-spring>

Lab 1: Deploy Zeebe Cluster & Console on Kubernetes

1. Clone Helm Chart Repo

```
git clone https://github.com/camunda/camunda-platform-deployment.git
```

```
cd camunda-platform-deployment/helm
```

...

2. Add Camunda Helm Repo & Update

```
helm repo add camunda https://helm.camunda.io
```

```
helm repo update
```

...

3. Create Namespace

```
kubectl create namespace camunda
```

...

4. Deploy Zeebe Cluster

```
helm install zeebe camunda/zeebe \
  --namespace camunda \
  --set broker.replicas=3 \
  --set gateway.replicas=2 \
  --set resources.broker.requests.cpu="500m" \
  --set resources.broker.requests.memory="1Gi" \
  --set resources.broker.limits.cpu="1" \
  --set resources.broker.limits.memory="2Gi"
...
```

5. Deploy Console (Operate, Tasklist, Optimize)

```
helm install console camunda/console \
  --namespace camunda \
  --set operate.enabled=true \
  --set tasklist.enabled=true \
  --set optimize.enabled=true
...
```

6. Verify Deployment

```
kubectl get pods -n camunda
...
```

All Zeebe broker and console pods should be in ****Running**** state.

Lab 2: Model & Deploy a Sample BPMN Process

1. Clone Sample Project

```
git clone https://github.com/camunda-community-hub/camunda-8-get-started-spring.git
cd camunda-8-get-started-spring
...
```

2. Open BPMN in Web Modeler

1. Navigate to `model/` folder
2. Open `order-process.bpmn` in Camunda Web Modeler (<https://camunda.com/web-modeler>)
3. Review the **Order Process** diagram: start event, service task, user task, end event

3. Export BPMN File

- In Web Modeler, click **Save** → **Download File** → save as `order-process.bpmn` in project root.

4. Deploy Process via Zeebe CLI

```
zbctl deploy order-process.bpmn
...
```

5. Instantiate Process

```
zbctl create instance OrderProcess
...
```

6. Access Tasklist

- Port-forward Tasklist: `kubectl port-forward svc/console-tasklist 8081:80 -n camunda`
- Open <http://localhost:8081> and complete the **Accept Order** user task.

Lab 3: External Elasticsearch Configuration

1. Deploy Elasticsearch via Helm

```
helm repo add elastic https://helm.elastic.co
```

```
helm repo update
```

```
helm install elasticsearch elastic/elasticsearch \
```

```
--namespace camunda \
```

```
--set replicas=3 \
```

```
--set resources.requests.cpu="500m" \
```

```
--set resources.requests.memory="1Gi" \
```

```
--set resources.limits.cpu="1" \
```

```
--set resources.limits.memory="2Gi"
```

...

2. Update Console Values

Create `external-es-values.yaml`:

```
operate:
```

```
  elasticsearch:
```

```
    addresses:
```

```
      - elasticsearch-master.camunda.svc.cluster.local:9200
```

```
tasklist:
```

```
  elasticsearch:
```

```
    addresses:
```

```
      - elasticsearch-master.camunda.svc.cluster.local:9200
```

...

3. Upgrade Console

```
helm upgrade console camunda/console \
```

```
--namespace camunda \
```

```
-f external-es-values.yaml
```

```
...
```

4. Validate

- In Operate UI (port-forward on port 8080), ensure you see process instances and job logs.

```
---
```

Lab 4: Integrate Keycloak for Cockpit Authentication

1. Deploy Keycloak

```
docker run -d --name keycloak -p 8081:8080 quay.io/keycloak/keycloak:latest start-dev
```

```
...
```

2. Configure Keycloak

1. Open <http://localhost:8081>

2. Login default user `admin`/`admin`

3. Create realm `camunda`

4. Create client `cockpit-client` with:

- Access Type: `confidential`

- Valid Redirect URI: `*`

5. Obtain **Client ID** and **Secret** from **Credentials** tab.

3. Update Helm Values

Create `keycloak-values.yaml`:

console:

authentication:

enabled: true

clientId: "cockpit-client"

clientSecret: "<SECRET_FROM_KEYCLOAK>"

issuerUri: "http://keycloak.local:8081/realms/camunda"

...

4. Redeploy Console

helm upgrade console camunda/console \

--namespace camunda \

-f keycloak-values.yaml

...

5. Verify

- Access Cockpit (port-forward on port 8082) and login via Keycloak.

Lab 5: Secrets Management with HashiCorp Vault

1. Deploy Vault

helm repo add hashicorp https://helm.releases.hashicorp.com

helm repo update

helm install vault hashicorp/vault \

--namespace camunda \

--set server.dev.enabled=true

...

2. Initialize & Unseal (Dev Mode skips unseal)

```
kubectl exec -n camunda vault-0 -- vault status
```

...

3. Store Broker Password in Vault

```
export VAULT_ADDR="http://$(kubectl get svc vault -n camunda -o  
jsonpath='{.status.loadBalancer.ingress[0].ip}'):8200"
```

```
vault kv put secret/camunda brokerPassword=mySecretPass
```

...

4. Deploy External Secrets Operator

```
kubectl apply -f https://raw.githubusercontent.com/external-secrets/kubernetes-external-  
secrets/main/deploy/crds.yaml
```

```
kubectl apply -f https://raw.githubusercontent.com/external-secrets/kubernetes-external-  
secrets/main/deploy/operator.yaml
```

...

5. Create ExternalSecret

Save `external-secret.yaml`:

```
apiVersion: external-secrets.io/v1beta1
```

```
kind: ExternalSecret
```

```
metadata:
```

```
  name: zeebe-broker-secret
```

```
  namespace: camunda
```

```
spec:
```

```
  refreshInterval: "1h"
```

```
  secretStoreRef:
```

```
    name: vault
```

```
kind: ClusterSecretStore
target:
  name: zeebe-broker-secret
  creationPolicy: Merge
data:
  - secretKey: brokerPassword
    remoteRef:
      key: secret/camunda
      property: brokerPassword
...

Apply:
kubectl apply -f external-secret.yaml
...
```

6. Update Zeebe Deployment to Use Secret

Set in Helm override `vault-secret-values.yaml`:

```
zeebe:
  broker:
    env:
      - name: ZEEBE_BROKER_SECURITY_SSL_KEY_PASSWORD
        valueFrom:
          secretKeyRef:
            name: zeebe-broker-secret
            key: brokerPassword
...
```

Upgrade:

```
helm upgrade zeebe camunda/zeebe -n camunda -f vault-secret-values.yaml
```

...

7. Validate

- Ensure Zeebe brokers start without errors and connect securely.