

1) Base OS prerequisites

Install base tools

```
apt-get install -y apt-transport-https ca-certificates curl gnupg lsb-release software-properties-common
```

Kernel modules and sysctls for Kubernetes networking

```
modprobe overlay
```

```
modprobe br_netfilter
```

```
tee /etc/sysctl.d/99-kubernetes-cri.conf >/dev/null <<'EOF'
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
sysctl --system
```

2) Install containerd (recommended for kubeadm)

Install containerd from Ubuntu repos (simple and stable)

```
apt-get install -y containerd
```

Generate default config and switch to systemd cgroups (important!)

```
mkdir -p /etc/containerd
```

```
containerd config default | tee /etc/containerd/config.toml >/dev/null
```

```
sed -i 's/SystemdCgroup = false/SystemdCgroup = true/' /etc/containerd/config.toml
```

Use systemd-resolved for DNS if needed (optional but helps pull issues on some hosts):

```
# sed -i 's|sandbox_image = ".*"|sandbox_image = "registry.k8s.io/pause:3.9"|' /etc/containerd/config.toml
```

```
systemctl daemon-reload
```

```
systemctl enable --now containerd
```

Note: We're not installing Docker Engine since Kubernetes uses containerd directly. If you still want Docker CLI for local builds, install it **after** Kubernetes is up (Section 8).

3) Install Kubernetes (kubeadm / kubelet / kubectl)

```
# Add Google Cloud apt repo for Kubernetes
```

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.30/deb/Release.key | gpg --dearmor -o  
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
```

```
https://pkgs.k8s.io/core:/stable:/v1.30/deb/ ' \
```

```
| tee /etc/apt/sources.list.d/kubernetes.list
```

```
apt-get update
```

```
# Install a consistent version set (example: 1.30.x)
```

```
apt-get install -y kubelet kubeadm kubectl
```

```
apt-mark hold kubelet kubeadm kubectl
```

```
# Ensure swap is off (again)
```

```
swapoff -a
```

4) Create the single-node cluster (kubeadm + Flannel)

We'll give the cluster a pod CIDR that Flannel expects: **10.244.0.0/16**.

```
# Pull control-plane images (optional, kubeadm will do it anyway)
```

```
kubeadm config images pull
```

```
# Initialize the cluster
```

```
kubeadm init --pod-network-cidr=10.244.0.0/16
```

Configure kubectl for your user

```
mkdir -p $HOME/.kube
```

```
cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
chown $(id -u):$(id -g) $HOME/.kube/config
```

Allow scheduling on the control-plane node (single-node setup)

```
kubectl taint nodes --all node-role.kubernetes.io/control-plane- || true
```

```
kubectl taint nodes --all node-role.kubernetes.io/master- || true
```

Install **Flannel** CNI (avoids your previous Tigera/Calico CRD annotation bloat):

```
kubectl apply -f https://raw.githubusercontent.com/flannel-  
io/flannel/master/Documentation/kube-flannel.yml
```

Wait for node to become Ready

```
kubectl get nodes -w
```

If the node stays NotReady, check:

```
kubectl -n kube-flannel get pods -o wide
```

```
kubectl describe node $(hostname)
```

```
journalctl -u kubelet -f
```

Common fixes: ensure sysctls in Section 1 are applied, containerd is running, and DNS resolves.

5) Install Helm (cleanly)

Using apt (or use official script if you prefer)

```
curl https://baltocdn.com/helm/signing.asc | gpg --dearmor -o  
/usr/share/keyrings/helm.gpg
```

```
echo "deb [signed-by=/usr/share/keyrings/helm.gpg]  
https://baltocdn.com/helm/stable/debian/ all main" \  
| tee /etc/apt/sources.list.d/helm-stable-debian.list
```

```
apt-get update
```

```
apt-get install -y helm
```

```
helm version
```

6) (Optional) Metrics & sanity checks

```
# Metrics Server (optional, nice for HPA & kubectl top)
```

```
kubectl apply -f https://github.com/kubernetes-sigs/metrics-  
server/releases/latest/download/components.yaml
```

```
# Watch core components
```

```
kubectl get pods -A -w
```

7) Install Camunda 8 with Helm (single-node friendly)

We'll put Camunda in its own namespace and use a **values.yaml** tuned for a single node.

Key points that address today's pain:

- Don't patch Elasticsearch anti-affinity live; configure it via Helm values.
- Keep replicas at 1.
- Keep resources light; disable Optimize if low on RAM.

```
helm repo add camunda https://helm.camunda.io
```

```
helm repo update
```

```
kubectl create namespace camunda
```

Create camunda-values.yaml:

```
cat > camunda-values.yaml <<'EOF'
```

```
# Minimal single-node friendly Camunda 8 (Camunda Platform)
```

```
# Disables strict anti-affinity so ES/Zeebe can schedule on the single node.
```

```
global:
```

```
image:
```

```
# You can pin versions if needed, otherwise defaults to chart defaults
```

tag: 8.6.0

ingress:

enabled: false # we'll use port-forward or NodePort

zeebe:

enabled: true

clusterSize: 1

partitionCount: 1

replicationFactor: 1

resources:

requests:

cpu: "200m"

memory: "512Mi"

limits:

cpu: "1"

memory: "2Gi"

elasticsearch:

enabled: true

master:

replicas: 1

persistence:

enabled: false

nodeSelector: {}

tolerations: []

affinity: {} # disables requiredDuring... anti-affinity that blocked you before

To reduce memory footprint in dev:

clusterHealthCheckParams: "wait_for_status=yellow&timeout=1s"

operate:

enabled: true

resources:

requests:

cpu: "100m"

memory: "256Mi"

limits:

cpu: "500m"

memory: "512Mi"

tasklist:

enabled: true

resources:

requests:

cpu: "100m"

memory: "256Mi"

limits:

cpu: "500m"

memory: "512Mi"

identity:

enabled: true

keycloak:

enabled: true

resources:

requests:

cpu: "100m"

memory: "256Mi"

limits:

cpu: "500m"

memory: "512Mi"

optimize:

enabled: false # disable to save memory in a small VM

EOF

Install:

helm install camunda camunda/camunda-platform -n camunda -f camunda-values.yaml

Watch pods come up (first pull can take a bit)

kubectl -n camunda get pods -w

When all are Running/Ready, you can port-forward to access the apps:

Operate (workflow management UI)

kubectl -n camunda port-forward svc/camunda-operate 8081:80

Tasklist

kubectl -n camunda port-forward svc/camunda-tasklist 8082:80

Identity/Keycloak login (Camunda web entry)

kubectl -n camunda port-forward svc/camunda-identity 8083:80

Open:

- <http://localhost:8083> (Identity/Keycloak)
Default credentials vary by chart; typically:
- **Keycloak** admin: admin / admin (or as set by chart)
- **Camunda** default users/clients are created by the chart; after logging into Identity, you can reach Operate and Tasklist via their URLs.

If you prefer NodePorts instead of port-forwards, set `service.type: NodePort` for each component in `camunda-values.yaml`.

8) (Optional) Install Docker Engine CLI (for local image builds)

Kubernetes already uses containerd. If you also want Docker CLI:

```
# Docker official repo
```

```
install -m 0755 -d /etc/apt/keyrings
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | \
```

```
  gpg --dearmor -o /etc/apt/keyrings/docker.gpg
```

```
chmod a+r /etc/apt/keyrings/docker.gpg
```

```
echo \
```

```
  "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] \
```

```
  https://download.docker.com/linux/ubuntu $(. /etc/os-release && echo  
  $VERSION_CODENAME) stable" \
```

```
  | tee /etc/apt/sources.list.d/docker.list > /dev/null
```

```
apt-get update
```

```
apt-get install -y docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-  
plugin
```

```
# Add your user to docker group (log out/in after)
```

```
usermod -aG docker $SUDO_USER 2>/dev/null || usermod -aG docker $USER
```

9) Health checks & common fixes (covers today's errors)

Node NotReady:

```
kubectl get nodes
```

```
kubectl describe node $(hostname)
```

```
journalctl -u kubelet -n 200 --no-pager
```

- Ensure containerd is running, swap is disabled, Section 1 sysctls applied, Flannel pods are healthy.

- If DNS issues cause image pull failures, verify `/etc/resolv.conf` points to `systemd-resolved` or a valid upstream; restart `containerd` after changes.

Image pull errors (ErrImagePull / ImagePullBackOff):

```
kubectl -n camunda describe pod <pod-name> | sed -n '/Events/, $p'
```

```
ctr -n k8s.io images list | head
```

- Retry: `kubectl delete pod <pod-name>` (Deployment will re-create).
- If a private registry is involved, create an `imagePullSecret` and reference it in values.

Elasticsearch anti-affinity / scheduling failures:

- We disabled strict anti-affinity in `camunda-values.yaml`. Don't live-patch the `StatefulSet`; update via Helm values and helm upgrade.

Calico/Tigera CRD annotation bloat error:

- We've avoided Calico entirely by using Flannel. If you must use Calico later, **install only the Tigera operator + a matching Calico manifest version** and avoid layering multiple kustomize bases that inflate annotations.

Helm chart upgrades / reconfig:

Edit values then:

```
helm upgrade camunda camunda/camunda-platform -n camunda -f camunda-values.yaml
```

If things get weird:

```
helm -n camunda get values camunda
```

```
helm -n camunda get manifest camunda | less
```

Clean uninstall of Camunda (namespace-scoped):

```
helm -n camunda uninstall camunda
```

```
kubectl delete ns camunda
```

10) Quick end-to-end verification

1) Cluster Ready

```
kubectl get nodes
```

2) CNI Ready

```
kubectl -n kube-flannel get pods
```

3) Camunda services exposed

```
kubectl -n camunda get svc
```

4) Pods healthy

```
kubectl -n camunda get pods
```

5) Login via Identity and reach Operate/Tasklist (via port-forward or NodePort)
